

Intro to Natural Language Processing



Kirill Kireyev
General Assembly

My Background

- PhD Computer / Cognitive Science
 - University of Colorado
- R&D @ Pearson Education
 - Adaptive vocabulary instruction
 - Essay scoring
 - Grammar/writing feedback

instaGrok

Grok

Journal

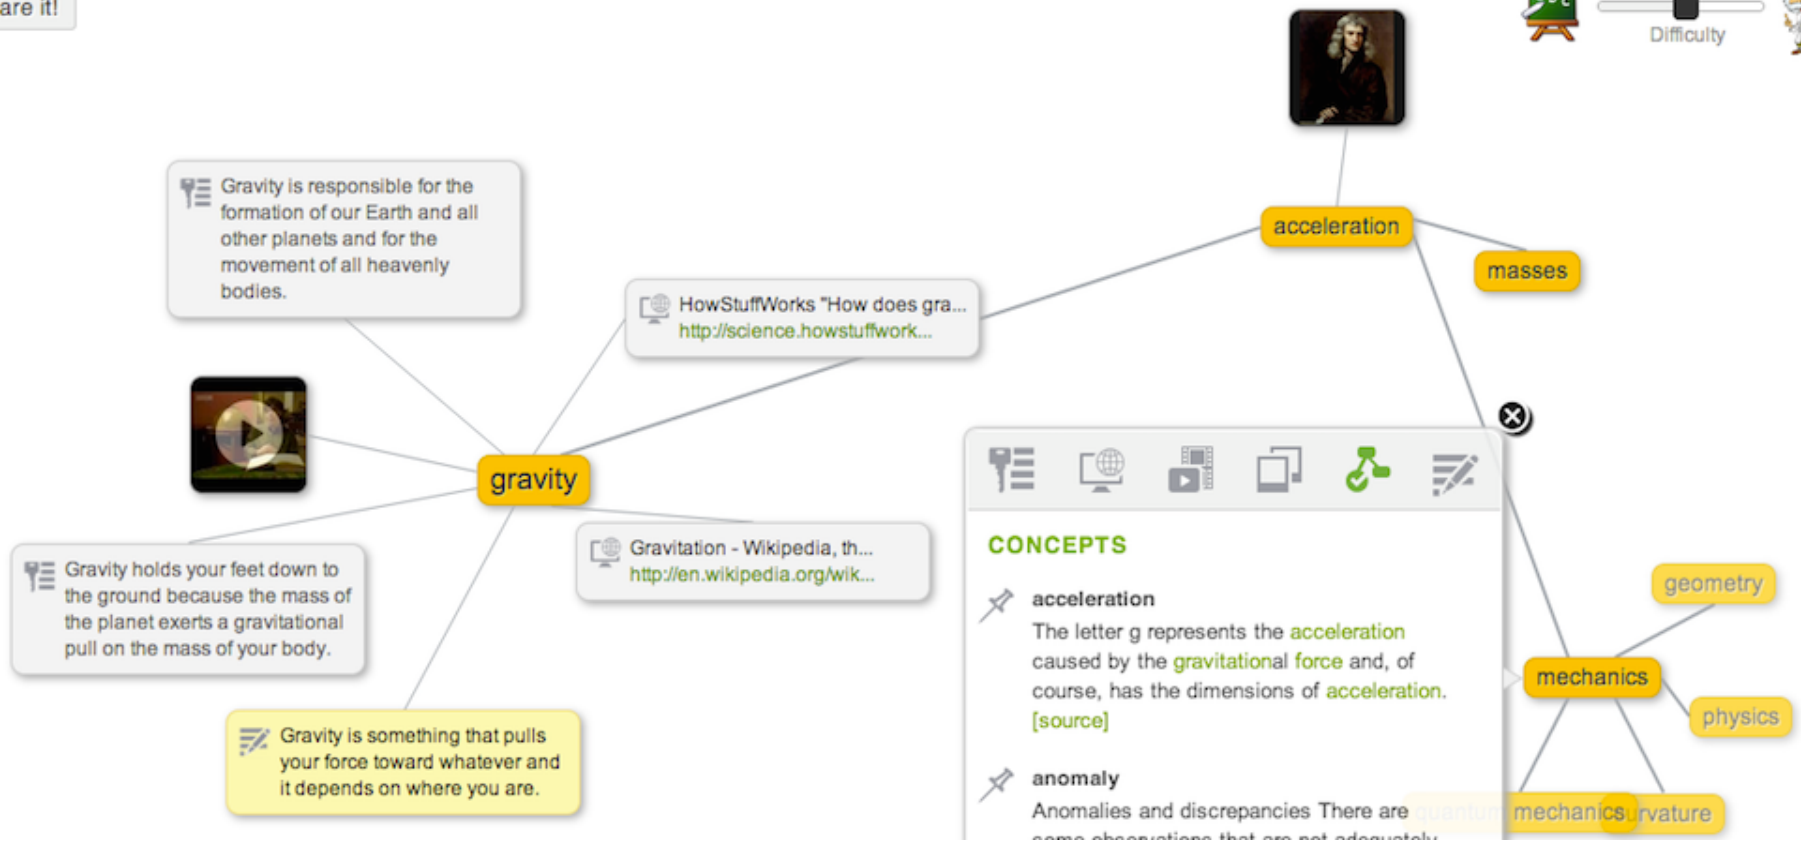
? Quizzes

Project: Gravity Project [edit]

Share it!



Difficulty



CONCEPTS

acceleration

The letter g represents the **acceleration** caused by the **gravitational force** and, of course, has the dimensions of **acceleration**.
[source]

anomaly

Anomalies and discrepancies There are **quantum mechanics** some observations that are not adequately

geometry

mechanics

physics

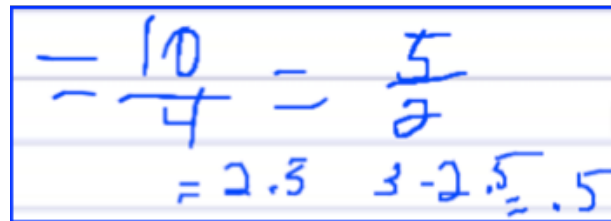
mechanics curvature

TEXT GENOME



Title	Length	Lexile	Grade	Vocab Words	Themes	Other
The Zoo Crew	582	763	5	healthy ranger volunteer	Animals (45) Types of people (14) Ownership/possession (13) Disease/health (11) Destructive/helpful (10)	<div> Rare words </div> <div> Age of Acq. </div> <div> Long Words </div> <div> Long Sents. </div> <div> Core Words </div> <div> Tech Words </div>

CATZILLA™


$$\begin{array}{r} 10 \\ - 4 \\ \hline 2.5 \end{array} \quad \begin{array}{r} 5 \\ - 2 \\ \hline 3 \\ - 2.5 \\ \hline 0.5 \end{array}$$

Optical Character Recognition (OCR),
Machine Learning (ML),
Natural Language Processing (NLP)

"Converts
to decimals"

ML,
rules

Recommendation

Explain how you know?

Because every time it is a new step it adds two more dots.

ML,
NLP

"Found
pattern"

ML,
rules

“Simple” Problem

Count the number of sentences in a text.

How would you do it?

Count the punctuation?

Regular expressions

```
In [1]: import re
```

```
In [2]: text = "My name is Kirill, I'm 36 years old."
```

```
In [3]: digits = re.compile('\d+')
```

```
In [4]: digits.search(text).group()
```

```
Out[4]: '36'
```

Sentence segmentation with regular expressions

```
In [1]: text = "My name is Kirill. I'm 36 years old. I like data science."
```

```
In [2]: import re
```

```
In [6]: punct = re.compile('[\.\?!\!]+')
```

```
In [9]: print len(punct.split(text))-1
```

3

Not so simple!

“The incident took place at 2 p.m. on Monday.”

“The federal government has begun a civil rights investigation into how Freddie Gray died of a spinal injury after he was arrested in Baltimore, the U.S. Justice Department said Tuesday.”

“Mayor Stephanie Rawlings-Blake and Police Commissioner Anthony W. Batts called for calm to allow police to complete their investigation.”

Using Machine Learning

- Annotate a lot of sentences
- Train a classifier
- Features?
 - Preceding word(s)
 - Following word(s)
 - Capitalization
 - Length of word

NLTK

```
In [1]: import nltk.data
```

```
In [2]: text = '''
... Punkt knows that the periods in Mr. Smith and Johann S. Bach
... do not mark sentence boundaries.  And sometimes sentences
... can start with non-capitalized words.  i is a good variable
... name.
... '''
```

```
In [3]: text
```

```
Out[3]: '\nPunkt knows that the periods in Mr. Smith and Johann S. Bach\ndo not mark sentence boundaries.  And sometimes sent
ences\ncan start with non-capitalized words.  i is a good variable\nname.\n'
```

```
In [4]: sent_detector = nltk.data.load('tokenizers/punkt/english.pickle')
```

```
In [5]: print('\n-----\n'.join(sent_detector.tokenize(text.strip())))
```

```
Punkt knows that the periods in Mr. Smith and Johann S. Bach
do not mark sentence boundaries.
-----
And sometimes sentences
can start with non-capitalized words.
-----
i is a good variable
name.
```

Understanding Language



Part of Speech Tagging

I went to the park with my friend.

PRP VBD TO DT NN IN PRP\$ NN

I went to park my car.

Will you friend me on Facebook?

NLTK POS Tagging

```
[1]: import nltk
```

```
[2]: tok = nltk.word_tokenize("I went to the park with my friend.")
```

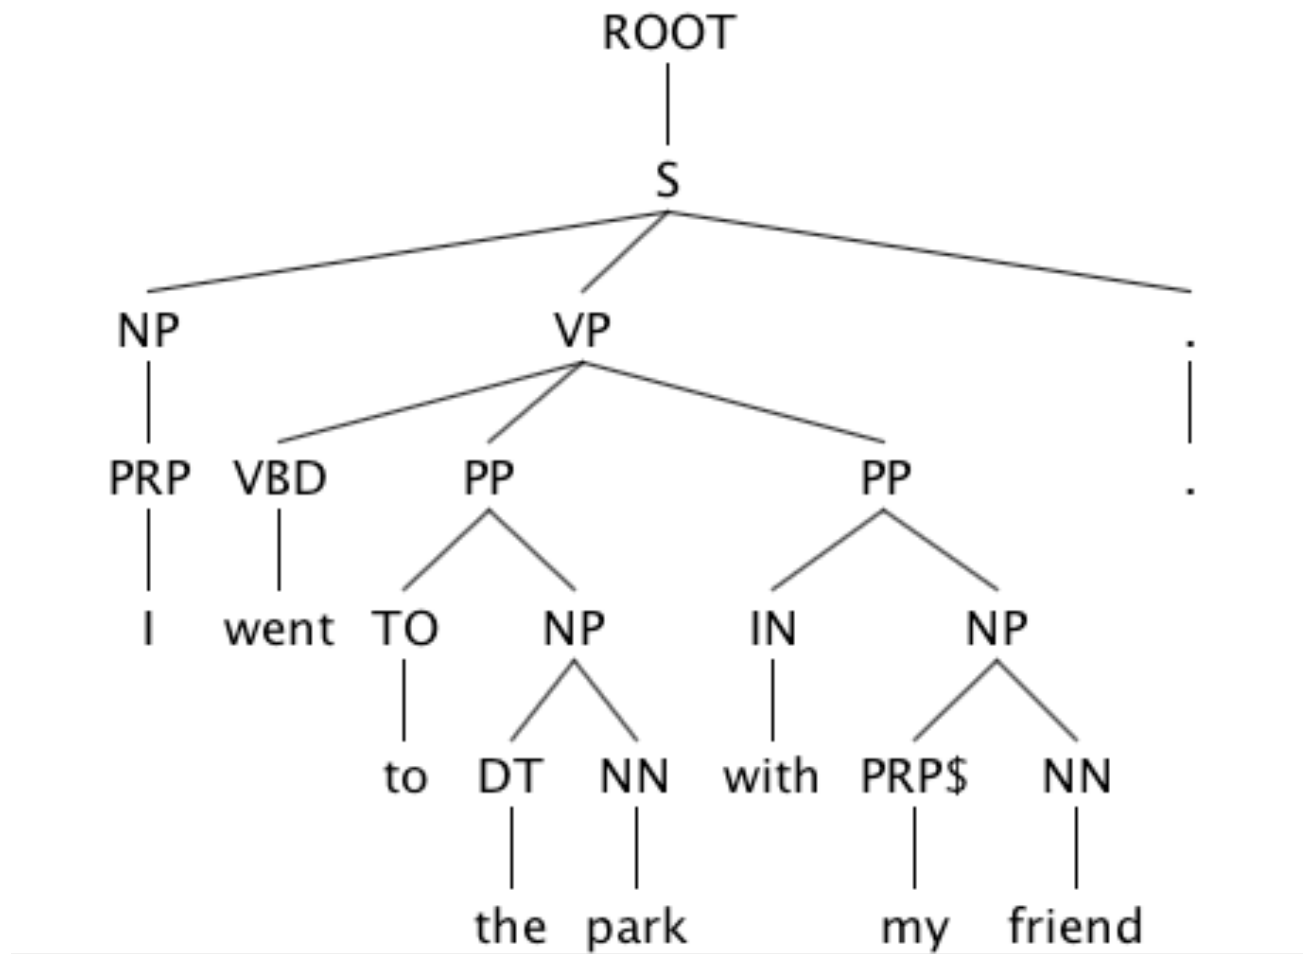
```
[3]: tok
```

```
t[3]: ['I', 'went', 'to', 'the', 'park', 'with', 'my', 'friend', '.']
```

```
[4]: nltk.pos_tag(tok)
```

```
t[4]: [('I', 'PRP'),  
      ('went', 'VBD'),  
      ('to', 'TO'),  
      ('the', 'DT'),  
      ('park', 'NN'),  
      ('with', 'IN'),  
      ('my', 'PRP$'),  
      ('friend', 'NN'),  
      ('.', '.')] 
```



Parsing



Parsing Ambiguity


While hunting in Africa, I shot an elephant in my pajamas.

How he got into my pajamas, I don't know.



Vector Representation

- Instead of deep semantic representation, represent a text with a vector (of words)
- “I went with my friend to the park” -> { “go”, “friend”, “park” }
- Remove *stopwords* (“to”)
- Dimensionality = { # words in lexicon }
 - Each word = independent dimension
- Ignores word order (“bag of words”)
- Can do similarity comparisons using linear algebra methods (e.g. cosine)
- Future: dimensionality reduction (SVD)



WordNet

- <http://wordnet.princeton.edu>
- Example: Finding synonyms NLTK Wordnet

```
In [1]: import nltk  
        from nltk.corpus import wordnet as wn
```

```
In [2]: wn.synsets('flabbergasted')
```

```
Out[2]: [Synset('flabbergast.v.01'), Synset('dumbfounded.s.01')]
```

```
In [5]: wn.synsets('flabbergasted')[0].definition()
```

```
Out[5]: u'overcome with amazement'
```

Sentiment Analysis

- “iPhone 6 is a **fantastic** product”
- “iPhone is **difficult** to use”
- Words with *sentiment polarity*:
 - Simple: fantastic: +1 / difficult: -1
- How to obtain polarity values?
 - Use annotated texts
 - Reviews
- Retain words (features) with high information gain:
 - E.g. Chi squared

Sentiment Analysis Complications

- “iPhone 6 is a **fantastic** product”
- “iPhone is **difficult** to use”
- “iPhone 6 is not **great**”
- “iPhone 6 is not necessarily the most **amazing** product I’ve used”
- Domain-dependent